

An Overview of the Rust Programming Language

Jim Royer

CIS 352

April 29, 2019



CIS 352 | Rust Overview | 1 / 1

References

- *The Rust Programming Language* by S. Klabnik and C. Nichols, 2018.
<https://doc.rust-lang.org/book/>
- *Rust*, N. Matsakis: <http://www.slideshare.net/nikomatsakis/guaranteeing-memory-safety-in-rust-39042975>
- *Rust: Unlocking Systems Programming*, A. Turon, <http://www.slideshare.net/InfoQ/rust-unlocking-systems-programming>
- *Rust's ownership and move semantics*, T. Ohzeki, <http://www.slideshare.net/saneyuki/rusts-ownership-and-move-semantics>
- *The Rust Programming Language*, A. Crichton, a Google Tech Talk,
<https://www.youtube.com/watch?v=d1uraoHM8Gg>
- *CIS 198: Rust Programming, University of Pennsylvania, Spring 2016*,
<http://cis198-2016s.github.io>
- *Programming Rust*, by Jim Blandy and Jason Orendorff, O'Reilly Media, 2017.
- <https://insights.stackoverflow.com/survey/2018/#technology-most-loved-dreaded-and-wanted-languages>

I shamelessly filch images and entire slides from these folks.

CIS 352

Rust Overview | 3 / 1

www.rust-lang.org



Install Learn Tools Governance Community Blog

Rust

GET STARTED

[Version 1.34.0](#)

Empowering everyone to build reliable and efficient software.

The Rust 2018 Edition is here!

Why Rust?

Performance

Rust is blazingly fast and memory-efficient: with no runtime or garbage collector, it can power performance-critical services, run on embedded devices, and easily integrate with other languages.

Reliability

Rust's rich type system and ownership model guarantee memory-safety and thread-safety – and enable you to eliminate many classes of bugs at compile-time.

Productivity

Rust has great documentation, a friendly compiler with useful error messages, and top-notch tooling – an integrated package manager and build tool, smart multi-editor support with auto-completion and type inspections, an auto-formatter, and more.

CIS 352 | Rust Overview | 2 / 1

Sample code: Factorial, 1

A recursive version of factorial

```
fn fact_recursive(n: u64) -> u64 {
    match n {
        0 => 1,
        _ => n * fact_recursive(n-1)
    }
}

fn main () {
    for i in 1..10 {
        println!("{}\t{}", i, fact_recursive(i));
    }
}
```

- `u64` ≡ the type of unsigned 64bit integers
- “`n: u64`” ≡ n is of type `u64`
- “`-> u64`” ≡ the fn returns a `u64`-value
- `match` ≡ Haskell's `case` (with pattern matching)
- `1..10` ≡ an iterator for the numbers 1 through 9 (*Yes, I mean 9.*)

CIS 352

Rust Overview | 4 / 1

Sample code: Factorial, 2

A iterative version of factorial

```
fn fact_iterative(n: u64) -> u64 {  
    let mut i      = 1u64;  
    let mut result = 1u64;  
    while i<=n {  
        result *= i;  
        i += 1;  
    }  
    return result;  
}  
  
fn main () {  
    for i in 1..10 {  
        println!("{}\t{}",  
                i, fact_iterative(i));  
    }  
}
```

- “`let mut i`” ≡ declares a *mutable* variable `i`.
- “`let j`” ≡ declares a *immutable* variable `i`.
- `1u64` ≡ `u64`-version of 1.
- The compiler figures out the types of `i` and `result`.

Sample code: Factorial, 3

A iterator version of factorial

```
fn fact_iterator(n: u64) -> u64 {  
    (1..n+1).fold(1, |p, m| p*m)  
}  
  
fn main () {  
    for i in 1..10 {  
        println!("{}\t{}",  
                i, fact_iterator(i));  
    }  
}
```

- `|p, m| p*m` ≡ $\lambda p, m. (p * m)$
- In fact, `fact_iterator` ≡ `foldl (\p m->p*m) 1 [1..n]`

Back to brag list from www.rust-lang.org, 1

- type inference
Rust's type system is a cousin of ML's and Haskell's
- pattern matching
as in ML and Haskell and Swift and ...
- trait-based generics
in place of OO-classes, Rust uses a version of Haskell's type-classes.
This is much less bureaucratic than the standard OO framework.

Back to brag list from www.rust-lang.org, 2

- zero-cost abstractions
- efficient C bindings
- minimal runtime

These are performance goals borrowed from C++.

Stroustrup on C++:

C++ implementations obey the *zero-overhead principle*:
What you don't use, you don't pay for.

And further:

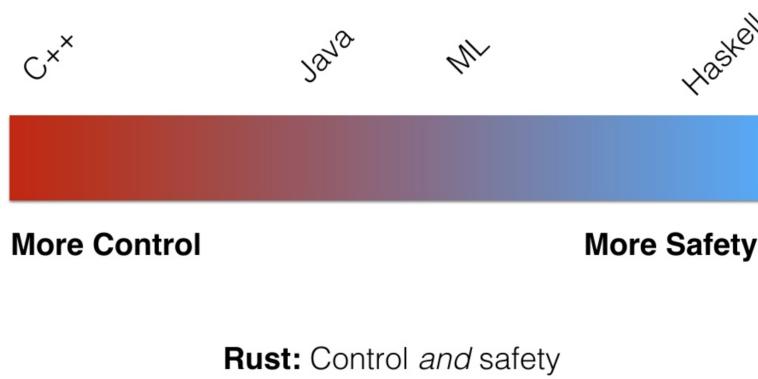
What you do use, you couldn't hand code any better.

- The prices Stroustrup is worried about paying are the *time and space* used by running programs.
- ... the time and effort binding wounds from repeated shooting yourself in the foot — that is *not* Stroustrup's concern.
- But Rust wants to avoid both sorts of prices.

Back to brag list from www.rust-lang.org, 3

- guaranteed memory safety
- threads without data races

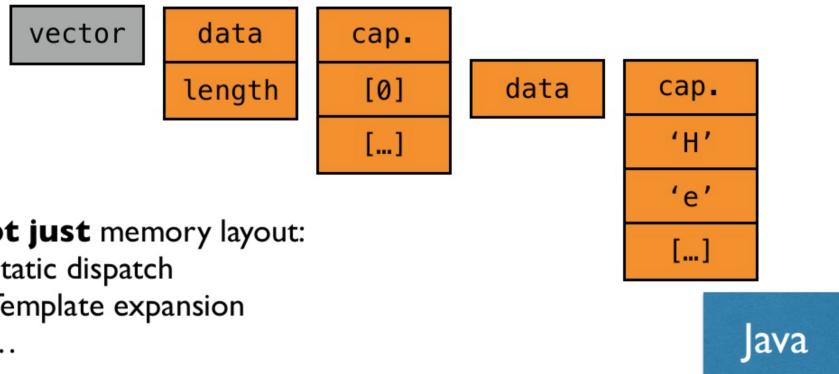
These are safety guarantees. Why this is a big deal.



CIS 352 Rust Overview 9 / 1

Zero-cost abstraction

Ability to define **abstractions** that **optimize away to nothing**.

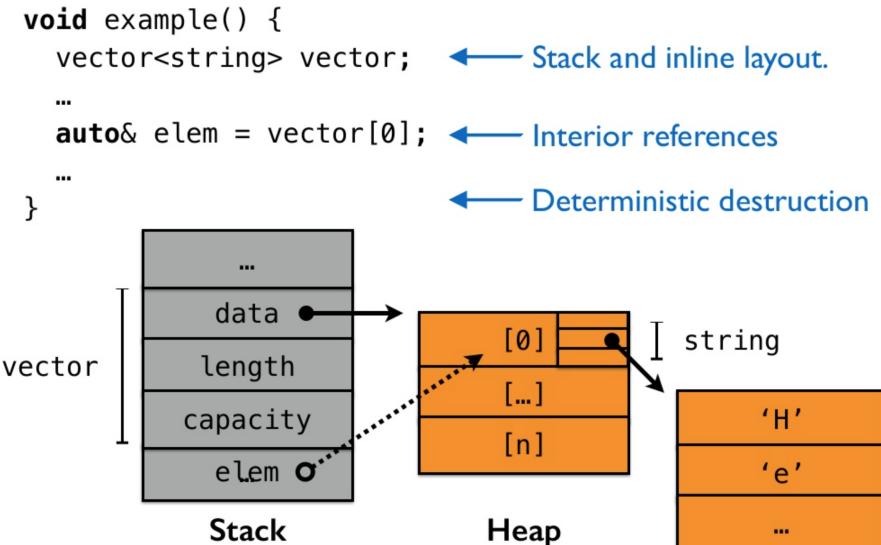


Not just memory layout:

- Static dispatch
- Template expansion
- ...

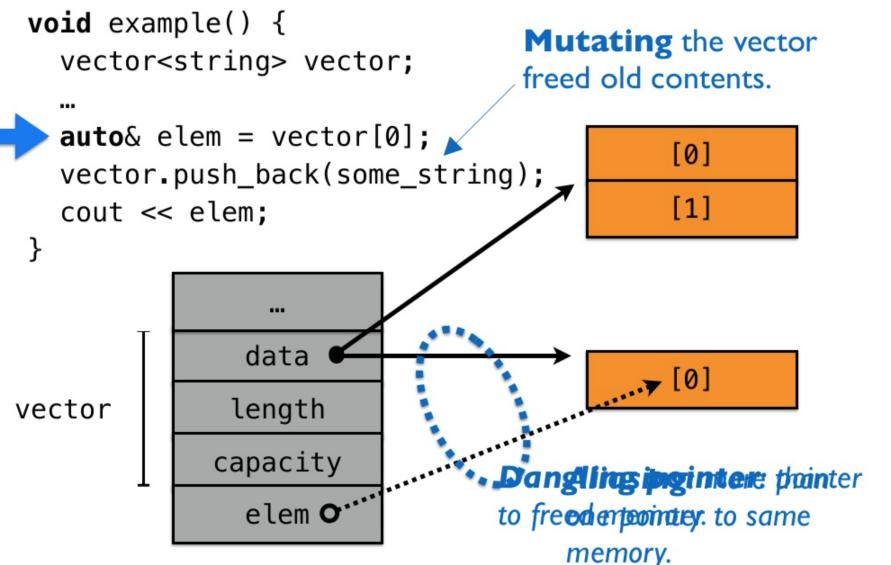
C++

What is **control**?



C++

What is **safety**?



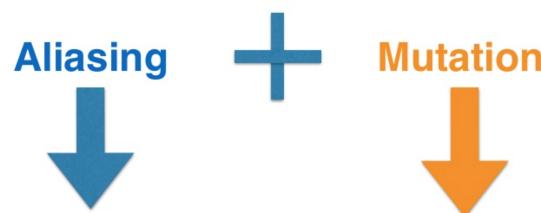
What about using a GC? (like Haskell, Java, Go, ...)

- Garbage collection:
 - The programmer allocates vectors, strings, etc.
 - The runtime system periodically sweeps through memory, looks for unreferenced data and deallocates it.
- ✗ Loss of control
- ✗ Runtime overhead
- ✗ Doesn't help with other safety issues: iterator invalidation, data races, etc.

So, what is Rust's solution? **Ownership**

Observation from C++-land

Danger arises from...



Hides dependencies.

Causes memory to be freed.

```
{ auto& e = v[0]; ... v.push_back(...); }
```

∴ Outlaw doing both at once and appoint the compiler Sheriff.

What about using a GC? (like Haskell, Java, Go, ...)

- Garbage collection:
 - The programmer allocates vectors, strings, etc.
 - The runtime system periodically sweeps through memory, looks for unreferenced data and deallocates it.
- ✗ Loss of control
- ✗ Runtime overhead
- ✗ Doesn't help with other safety issues: iterator invalidation, data races, etc.

So, what is Rust's solution? **Ownership** (*affine linear typing*)



Three Basic Patterns

Ownership	<code>fn foo(v: T) { ... }</code>
Shared Borrow	<code>fn foo(v: &T) { ... }</code>
Mutable Borrow	<code>fn foo(v: &mut T) { ... }</code>

The next bunch of slides are from
CIS 198: *Rust Programming, University of Pennsylvania, Spring 2016*,
<http://cis198-2016s.github.io/slides/01/>

Ownership

- A variable binding takes ownership of its data.
A piece of data can only have one owner at a time.
- When a binding goes out of scope, the bound data is released automatically.
For heap-allocated data, this means de-allocation.
- Data must be guaranteed to outlive its references.

```
fn foo() {  
    // Creates a Vec object.  
    // Gives ownership of the Vec object to v1.  
    let mut v1 = vec![1, 2, 3];  
    v1.pop();  
    v1.push(4);  
    // At the end of the scope, v1 goes out of scope.  
    // v1 still owns the Vec object, so it can be cleaned up.  
}
```

Move Semantics

```
let v1 = vec![1, 2, 3];  
let v2 = v1; // Ownership of the Vec object moves to v2.  
println!("", v1[2]); // error: use of moved value 'v1'
```

- `let v2 = v1;`
 - We don't want to copy the data, since that's expensive.
 - The data cannot have multiple owners.
 - **Solution:** move the Vec's ownership into v2, and declare v1 invalid.
- `println!("{}", v1[2]);`
 - We know that v1 is no longer a valid variable binding, ∵ error!
 - Rust can reason about this at compile time, ∵ compiler error.
 - Moving ownership is a compile-time semantic.
It doesn't involve moving data during your program.

Ownership does not always have to be moved

- Rust would be a pain to write if we were forced to explicitly move ownership back and forth.

```
fn vector_length(v: Vec<i32>) -> Vec<i32> {  
    // Do whatever here,  
    // then return ownership of 'v' back to the caller  
}
```

- The more variables you had to hand back (think 5+), the longer your return type would be!

Borrowing

- In place of transferring ownership, we can borrow data.
- A variable's data can be borrowed by taking a reference to the variable (i.e., aliasing); ownership doesn't change.
- When a reference goes out of scope, the borrow is over.
- The original variable retains ownership throughout.

```
let v = vec![1, 2, 3];  
let v_ref = &v; // v_ref is a reference to v.  
assert_eq!(v[1], v_ref[1]); // use v_ref to access the data  
                           // in the vector v.  
// BUT!  
let v_new = v; // Error, cannot transfer ownership  
               // while references exist to it.
```

Borrowing

```
// 'length' only needs 'vector' temporarily, so it is borrowed.
fn length(vec_ref: &Vec<i32>) -> usize {
    // vec_ref is auto-dereferenced when you call methods on it.
    vec_ref.len()
}
fn main() {
    let vector = vec![];
    length(&vector);
    println!("{}: {:?}", vector); // this is fine
}
```

- References, like bindings, are immutable by default.
- The borrow is over after the reference goes out of scope (at the end of `length`).
- (`usize` = The pointer-sized unsigned integer type.)

CIS 352

Rust Overview 20 / 1

Borrowing Rules

- ① You can't keep borrowing something after it stops existing.
- ② One object may have many immutable references to it (`&T`).
- ③ OR exactly one mutable reference (`&mut T`) (not both).

CIS 352

Rust Overview 22 / 1

Borrowing

```
// 'push' needs to modify 'vector' so it is borrowed mutably.
fn push(vec_ref: &mut Vec<i32>, x: i32) {
    vec_ref.push(x);
}
fn main() {
    let mut vector: Vec<i32> = vec![];
    let vector_ref: &mut Vec<i32> = &mut vector;
    push(vector_ref, 4);
}
```

- Variables can be borrowed by mutable reference: `&mut vec_ref`.
 - `vec_ref` is a reference to a mutable `Vec`.
 - The type is `&mut Vec<i32>`, not `&Vec<i32>`.
- Different from a reference which is variable.
- You can have exactly one mutable borrow at a time.
Also you cannot dereference borrows (changes ownership).

CIS 352

Rust Overview 21 / 1

Borrowing Prevents: Use-after-free bugs

Valid in C, C++,...

```
let y: &i32;
{
    let x = 5;
    y = &x; // error: 'x' does not live long enough
}
println!("{}: {}", *y);
```

This eliminates vast numbers of memory safety bugs *at compile time!*

CIS 352

Rust Overview 23 / 1

The ownership rules also prevent other things

Under the standard ownership rules:

- You **cannot** implement doubly-linked lists (and circular structures in general).
- You **cannot** call C libraries.
- ...

Unsafe Rust

```
unsafe {  
    ...  
}
```

- Relaxes some of the checking rules.
- Allows C libraries calls.
- Allows access to raw pointers.
- Allows you to implement language extensions, e.g., doubly-linked lists, garbage-collected pointers, etc.

CIS 352

Rust Overview 24 / 1

Concurrency

- The ownership rules also turn out to be useful for implementing safe concurrency (e.g., threads, interprocess communication, etc.)
- Standard currency primitives are not built in to Rust — they can be defined via ownership rules.
- Easy to use, safe (e.g., data-race free) concurrent programming is a big deal.

CIS 352

Rust Overview 25 / 1

Other goodies

- Crate and cargo — a modern, simple to use project manager to track dependencies, etc.
- Hygienic macros.
- Growing collection of libraries (nowhere as near mature or complete as C, C++, Java, etc.)
- Etc.
- See <https://www.rust-lang.org> for other resources.
- If you want to learn Rust, don't trust anything over three years old. (*Earlier versions were quite different.*)

CIS 352

Rust Overview 26 / 1

Nota Bene

- There isn't a type-soundness proof for Rust yet. That is, the story Rust tells about what its types mean is nice — but can you prove that the story is correct?
- **Problems:**
 - Rust is a moving target.
 - Rust includes unsafe-blocks.

CIS 352

Rust Overview 27 / 1